

## Unit 5 Sample Problems - Processes

In this sample problem set, we will practice concepts of processes and process management.

- Length: 60 minutes with discussion.
- Questions: Q1, Q3-Q5. (optional: Q2, Q6)

### Process Concept

1. [Acuña] There are four general regions to a process in memory: stack, heap, data, and text. **List and explain** which of these must be separated for each program running and which might be combined globally. (Don't worry about security or programs misbehaving.) [2 points]

Ans: [Acuña]

- Stack: This must be kept separate since programs must track their own execution.
- Heap: This can safely be combined - all programs are simply linking to allocations at different addresses.
- Data: This can safely be combined - variables will get stored in different places.
- Text: This can be problematic - code will get stored in different places but we will need some way to distinguish between different entry points (main functions).

2. [Bush] A PCB typically contains a list of the files that the process has open. **Explain** how the PCB concept could be used to avoid concurrent IO issues? [2 points]

Ans: [Bush]

Before the operating system allows a process to open a file for writing, it can check through the other PCBs to make sure that another process is not already writing to that file. The operating system could also prevent reads from files that are currently being written. (It should be noted that operating systems do not commonly prevent concurrent file access issues).

### Process Scheduling

3. [Bahremand] **Explain** why the queue diagram in slide 9 has a *continuous cycle* that flows between the listed queues and resources? Is a cycle necessary? [2 points]

Ans: [Bahremand]

All processes start in the ready queue, and then move to the CPU for computation before moving into another queue that requests resources such as I/O, forking of processes, and listening for interrupts to invoke other actions. Each of these resources are limited and are made available via a queue. Once the resource has been used, the process returns to the ready queue to determine if they can be terminated, or need another round of resource allocation. The ready queue is continuously looping through processes, and feeding it into the CPU to ensure that each process gets resources for as long as it needs while no particular one prevents others from getting CPU/IO time.

## Operations on Processes

4. [Lisonbee] **Trace** the program below, identify the values of the pids at lines A, B, C, D, E, and F. (Assume that the actual pid of Process 1 is 1502, Process 2 is 1505, and Process 3 is 1507. Also assume that fork will always succeed.) [4 points]

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main() {
    pid_t temp1, temp2, temp3;
    temp1 = fork();

    if (temp1 < 0) { /* Error occurred */
        fprintf(stderr, "Fork_Failed");
        return 1;
    }
    else if (temp1 == 0) { /* Process 2 */
        temp2 = fork();

        if (temp2 < 0) { /* Error occurred */
            fprintf(stderr, "Fork_Failed");
            return 1;
        }
        else if (temp2 == 0) { /* Process 3 */
            temp3 = getpid();
            printf("temp2=_%d", temp2); /* A */
            printf("temp3=_%d", temp3); /* B */
        }
        else { /* Process 2 */
            temp3 = getpid();
            printf("temp2=_%d", temp2); /* C */
            printf("temp3=_%d", temp3); /* D */
            wait(NULL);
        }
    }
    else { /* Process 1 */
        temp2 = getpid();
        printf("temp1=_%d", temp1); /* E */
        printf("temp2=_%d", temp2); /* F */
        wait(NULL);
    }
    return 0;
}
```

**Ans:** [Lisonbee]

A-0, B-1507, C-1507, D-1505, E-1505, F-1502

5. [Lisonbee] **Trace** the following program and then list all of the possible outputs that could be generated. (Assume that fork will always succeed.) [2 points]

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main() {
    pid_t pid1, pid2;
    pid1 = fork();

    if (pid1 < 0) {
        fprintf(stderr, "Fork_Failed");
        return 1;
    }
    else if (pid1 == 0) {
        printf("A");
        pid2 = fork();

        if (pid2 < 0) {
            fprintf(stderr, "Fork_Failed");
            return 1;
        }
        else if (pid2 == 0) {
            printf("B");
        }
        else {
            wait(NULL);
            printf("C");
        }
    }
    else {
        printf("D");
    }
    return 0;
}
```

**Ans:** [Lisonbee]

"ABCD", "ADBC", "DABC", "ABDC" (A always prints before B and C, and B always prints before C due to the wait(NULL) function call. D can print anytime.)

## Interprocess Communication

6. [Lisonbee] Consider a system where two processes (a producer and a consumer) use a message-passing system to communicate, and each process does work at different rates. The producer can produce (perform work) at any rate, but the consumer has to wait for the producer to complete its task

before moving on. Based on this system's needs, **explain** whether a synchronous or asynchronous communication system would be a better choice and why. [2 points]

**Ans:** [Lisonbee]

An asynchronous system would be better suited to this situation because information only flows in one direction (from producer to consumer). The producer can keep doing work even if the consumer isn't ready for it yet (as it can just be stored in a buffer for later use). Using a synchronous solution would limit the speed of the entire system to the slowest of the two processes.

## Examples of IPC Systems

### Communication in Client-Server Systems