

PROCESSES

Ruben Acuña

Spring 2018

Kernel source code is coming
from Linux 4.9.6. See
<https://www.kernel.org/>





PROCESS CONCEPT (.1)

We only have one \$sp register – how do we support multiple processes?

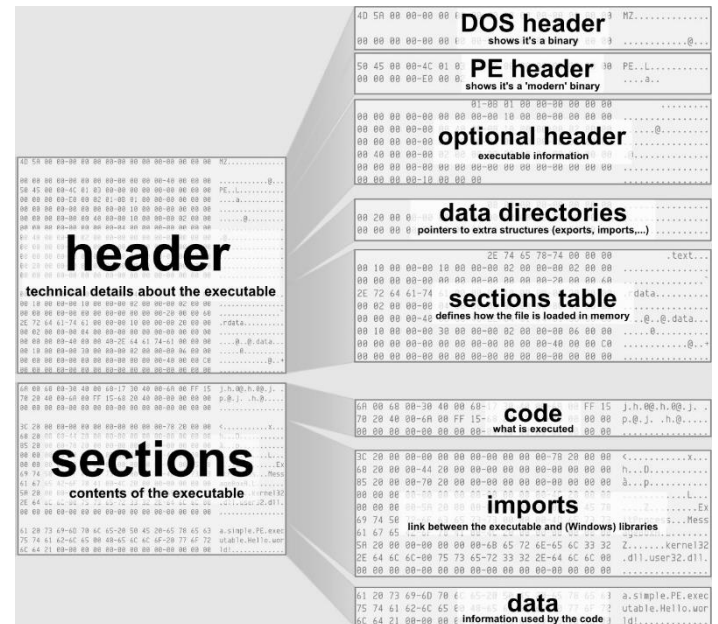
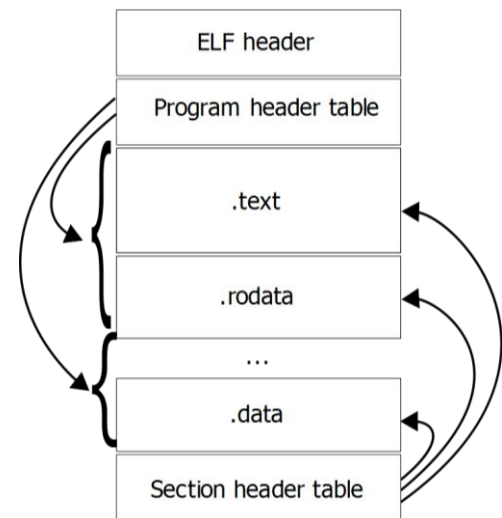
PROCESS AND STATE

- Roughly, a process is to an executable what an object is to a class.
 - Book also mentions the term *job*...
- An executing process has storage for stack, heap, “data”, and “text”.
 - Really: picture shows a “virtualized environment” that program is supposed to execute within.
- A general abstraction for the state of a process is shown to the right. In general, most operating systems will have these states although they may be named differently, or some may be combined/split. (Linux example to come.)



EXECUTABLE FORMATS

- All executables are obligated to follow a specific *Application Binary Interface* (ABI) that specifies system properties like type size, and library interface mechanism.
- Executable and Linkable Format (ELF; Unix)
 - Includes header, program, and various other sections (data).
- Portable Executable (PE; Win32)
 - Includes legacy MS DOS header, followed by modern PE header, and associated sections (text, data, rdata).



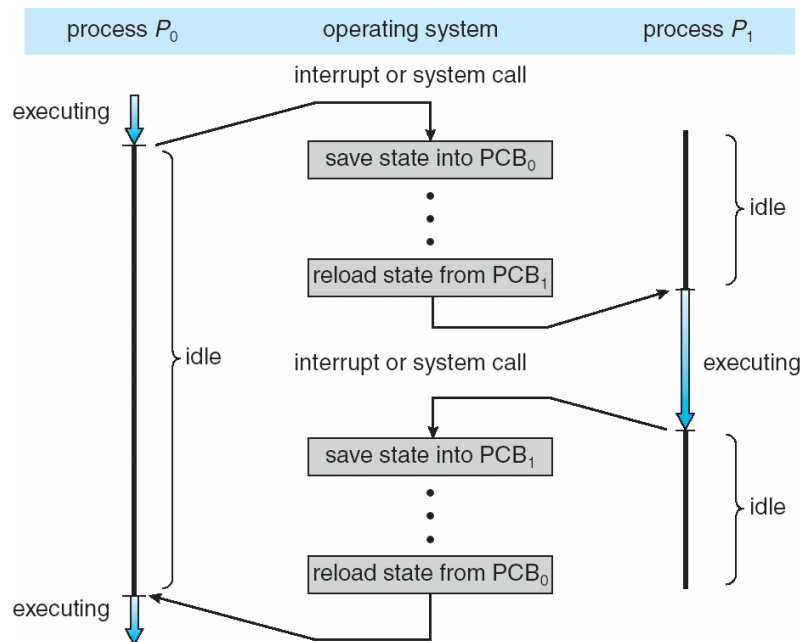
PROCESS CONTROL BLOCK

- Recall that processes live in a “virtualized environment” – they act as if they the only thing running on the CPU.

- We also need the ability to start and stop processes at will.

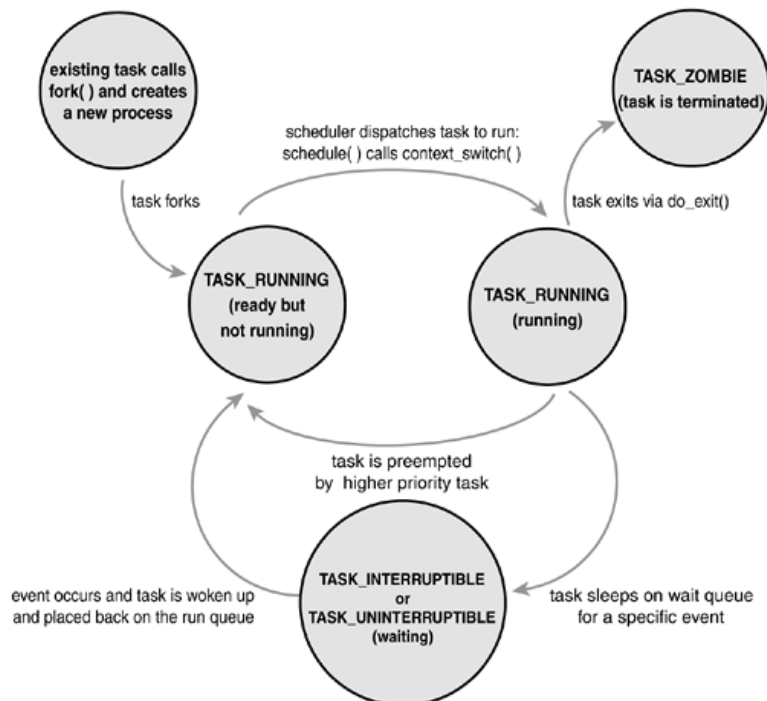
- Ergo, the entire state of a process is captured with something called an Process Control Block (PCB), which contains:

- Process state
- Program counter
- CPU registers
- CPU-scheduling information
- Memory management information
- Accounting information
- I/O status information



STRUCT TASK_STRUCT;

- In Linux, processes are represented with a *process descriptor* struct.
- Process descriptors are stored in a circular doubly linked list.



```
struct task_struct {
    volatile long state;           /* -1 unrunnable, 0 runnable, >0 stopped */
    void *stack;
    unsigned int flags;           /* per process flags, defined below */
    unsigned int ptrace;
```

```
int prio, static_prio, normal_prio;
unsigned int rt_priority;
const struct sched_class *sched_class;
```

```
unsigned int policy;
int nr_cpus_allowed;
cpumask_t cpus_allowed;
```

```
/* task state */
int exit_state;
int exit_code, exit_signal;
int pdeath_signal; /* The signal sent when the parent dies */
unsigned long jobctl; /* JOBCTL_*, siglock protected */
```

```
pid_t pid;
pid_t tgid;
```

```
struct task_struct __rcu *real_parent; /* real parent process */
struct task_struct __rcu *parent; /* recipient SIGCHLD, wait4() reports */
struct list_head children; /* list of my children */
struct list_head sibling; /* linkage in my parent's children list */
struct task_struct *group_leader; /* threadgroup leader */
```

```
/* RA: ~400 lines omitted. Mostly preprocessor directives. */
```

```
struct thread_struct thread;
};
```

TASK_STRUCT MEMORY ALLOCATION

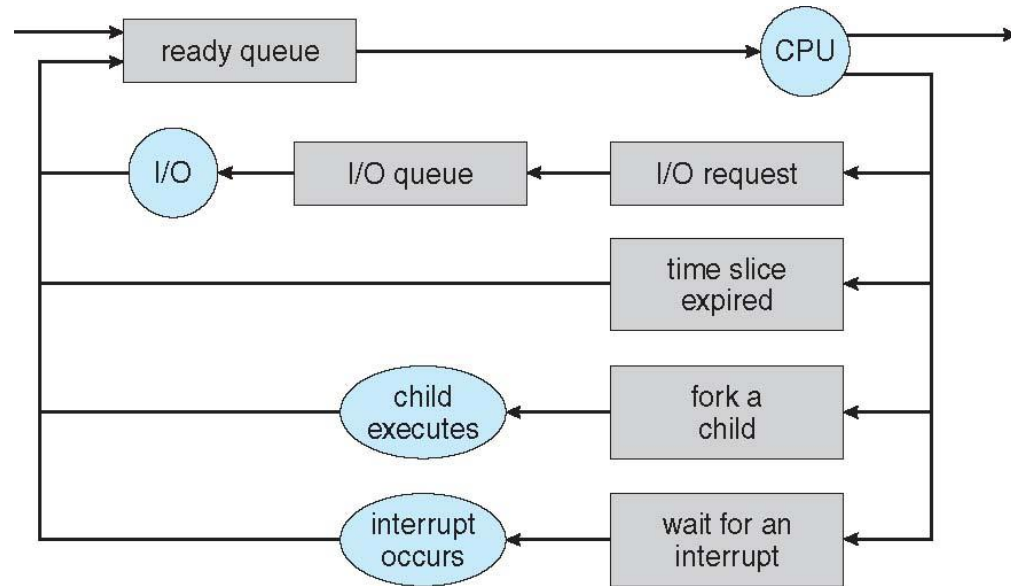
- task_structs come and go quickly, but memory management is slow - how to fix?



PROCESS SCHEDULING (.2)

SCHEDULING QUEUES

- Given that we want to support executing more than one process at a time, OSs need some sort of *process scheduler* that divides up resources (CPU time).
- We have two general notations:
 - *Ready queue*: processes ready to execute on the CPU.
 - *Device queue*: processes waiting to use a device/resource.
- Whenever a process from the ready queue is given CPU time, we say that it is *dispatched*.

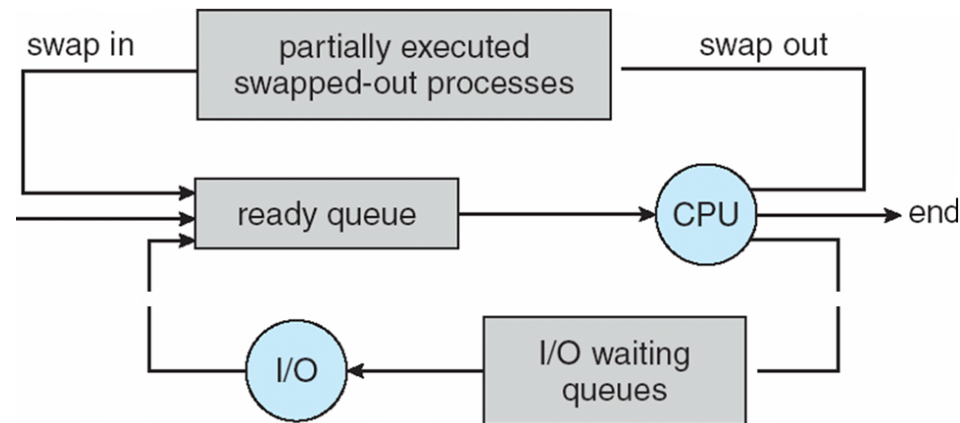


Queueing Diagram – rectangles indicates queues, ovals indicate resources.

Would it make sense to use priority queues?

SCHEDULERS

- There 2 patterns of process resource utilization:
 - IO-bound
 - CPU-bound
 - If we want to maximum performance, what ratio of these should we have and why?



Adding *swapping* into queueing system.

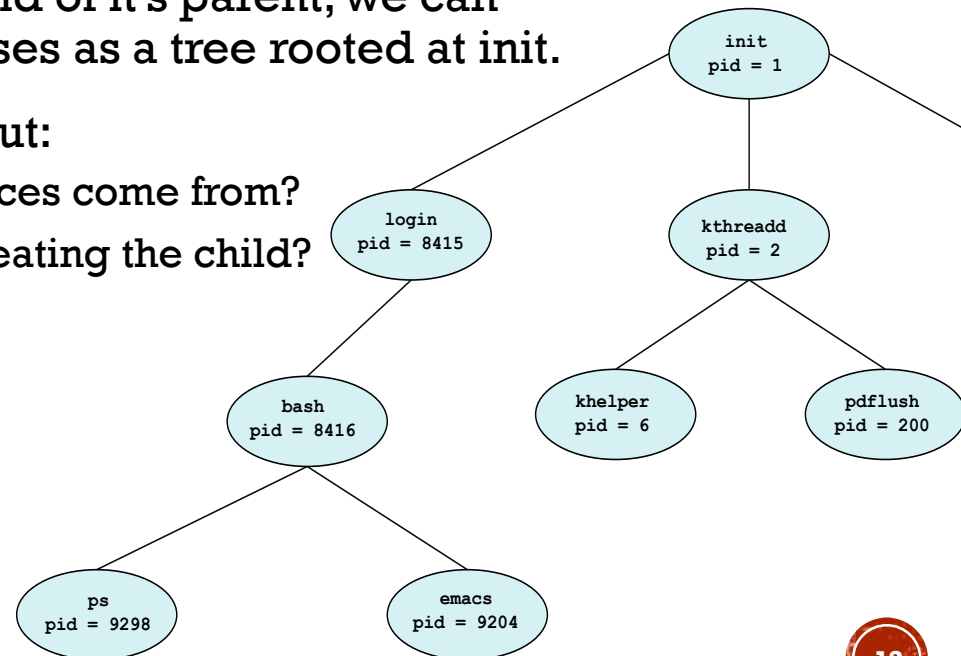
- There are 3 types of schedulers:
 - Short-term scheduler (RAM)
 - Medium-term scheduler (swap)
 - Long-term scheduler (Disk)
- During process execution, we say that the processes' *context* is active.
- If we want to change the running process, then we must preform a *context switch* to restore a PCB's state. We must do a state save as well.



OPERATIONS ON PROCESSES (.3)

CREATION

- As seen in `task_struct`, each process is identified by a `pid`.
 - In general, $pid2 > pid1$ which indicates that `pid2` was created later.
- The first process to be executed is called `init` and has `pid=1`. Since each process contains the `pid` of its parent, we can view the relation between processes as a tree rooted at `init`.
- We have some things to think about:
 - Where should the children's resources come from?
 - What should the parent do after creating the child?
 - What code should the child use?



POSIX SAMPLE — SPLITTING EXECUTION

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main() {
    pid_t pid;
    pid = fork();

    if (pid < 0) {
        //What process does this execute in?
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) {
        //What process does this execute in?
        execlp("/bin/ls", "ls", NULL);
    }
    else {
        //What process does this execute in?
        wait(NULL);
        printf("Child Complete");
    }
    return 0;
}
```



Under POSIX, fork() results in a child process that take is copied from the parent. Does copying seem like a good idea?

WINAPI - EXECUTING COMMAND

```
#include <stdio.h>
#include <windows.h>
int main(VOID) {
    STARTUPINFO si;
    PROCESS_INFORMATION pi;
    ZeroMemory(&si, sizeof(si));           /* RA: DOES NOT allocate memory */
    si.cb = sizeof(si);
    ZeroMemory(&pi, sizeof(pi));
    if (!CreateProcess(NULL,                /* lpApplicationName */
        "C:\\WINDOWS\\system32\\mspaint.exe", /* lpCommandLine */
        NULL,                                /* lpProcessAttributes */
        NULL,                                /* lpThreadAttributes */
        FALSE,                               /* bInheritHandles */
        0,                                   /* dwCreationFlags */
        NULL,                                /* lpEnvironment */
        NULL,                                /* lpCurrentDirectory */
        &si,
        &pi)) {
        fprintf(stderr, "Create Process Failed");
        return -1;
    }
    WaitForSingleObject(pi.hProcess, INFINITE);
    printf("Child Complete");
    CloseHandle(pi.hProcess);
    CloseHandle(pi.hThread);
}
```

IMPLEMENTING FORK()

- Here's a thought: isn't it expensive for `fork()` to create a copy of the original process if we immediately `exec()` over it?
- So, why not just omit copying all of the process state and instead use a reference pointer to it? Only when the execution of the forked process requires a write, then, perform the copy.
- This technique is called *copy-on-write* and can be seen as a *lazy* method of implementing *deep clone* operation on a data structure.



TERMINATION

- To terminate a process, one invokes `exit(n)`;
 - This is the implicit result of the return statement in main.
 - Or, may be invoked by an another program (parent, system, user, etc.).
- What are the reasons a parent might terminate a child?
- If a parent terminates, then typically each descendent process is also terminated (*cascading termination*). (Is this a good idea?)
- Under POSIX, we can use `wait(&n)` to “wait” for a child process to terminate. (Means we must save the pid from fork.)
- If a process terminates before it's parent has called `wait()`, then it cannot fully be removed from the process tree. (Why?) It is a called a *zombie* process. If somehow the parent is lost, then the executing child process may become *orphaned*.



INTERPROCESS COMMUNICATION (.4)

COMMUNICATION TYPES

- The process model we have seen so far is for *independent* processes, i.e., those which do not share data. The alternative is *cooperative* processes. There are reasons for cooperation:
 - Information Sharing
 - Computation Speedup
 - Modularity
 - Convenience
- Recall earlier, the two mechanisms of IPC we discussed: Shared Memory and Message-Passing.
 - Are there times when one in particular is necessary?
 - Which is faster in terms of system calls?

Note that this is all conceptual! The APIs are different.

SHARED-MEMORY SYSTEMS

- Producer-Consumer Problem:
 - Some process P1 creates products that are to be consumed by process P2.

- Let's try solving this problem with a buffered shared-memory solution.

- Does this code seem safe?

```
//shared data
#define BUFFER_SIZE 10
typedef struct {
    //...
} item;

item buffer[BUFFER_SIZE];
int in = 0, out = 0;

//Producer
item next_produced;
while (true) {
    /* produce next_produced item */
    while (((in + 1) % BUFFER_SIZE) == out);
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER_SIZE;
}

//Consumer
item next_consumed;
while (true) {
    while (in == out);
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    /* consume next_consumed item*/
}
```


MESSAGE-PASSING SYSTEMS

- At minimum, we need functions `send(msg)` and `receive(msg)`.
 - How much does it matter if the message is fixed length or variable length?
- If processes P1 and P2 want to communicate, we must establish a *communication link*.
- Processes may use *direct communication*. Links are complete, pairwise, and unique.
 - Symmetric: `send(P1, message)`, `receive(P2, message)`
 - Asymmetric: `send(P, message)`, `receive(&id, message)`
- Processes may also use *indirect communication over a mailbox/port*. Links are specific, groupwise, and non-unique.
 - `send(M, message)`, `receive(M, message)`
 - Consider the following problem: P1 sends a message to port K, but then P2 and P3 do a receive on K. What should happen? Is there an issue to fix?


MESSAGE-PASSING SYSTEMS

- For messaging commands we have a choice of either waiting for an answer (Blocking/Synchronous), or just continuing on (Nonblocking/Asynchronous):
 - Blocking Send
 - Non-blocking Send
 - Blocking Receive
 - Non-blocking receive
- By picking the correct operations, can we easily “solve” the producer-consumer problem...?
- There are three ways we can design a buffer:
 - Zero Capacity
 - Bounded Capacity
 - Unbounded Capacity

Blocking



Non-blocking



MESSAGE-PASSING SYSTEMS

- For messaging commands we have a choice of either waiting for an answer (Blocking/Synchronous), or just continuing on (Nonblocking/Asynchronous):

- Blocking Send
- Non-blocking Send
- Blocking Receive
- Non-blocking receive

- By picking the correct operations, can we easily “solve” the producer-consumer problem...?

- There are three ways we can design a buffer:

- Zero Capacity
- Bounded Capacity
- Unbounded Capacity

```
//producer
while (true) {
    /* produce next_produced item */
    send(next_produced);
}

//consumer
while (true) {
    receive(next_consumed);

    /* consume next_consumed item */
}
```



EXAMPLES OF IPC SYSTEMS (.5)

POSIX — SHARED MEMORY

- First, let's look at shared memory in POSIX. It functions via the idea of a memory-mapped file that one process creates and the other accesses.
- The next slide has a full program but four functions in particular are key:
 - `shm_open()`
 - `ftruncate()`
 - `mmap()`
 - `shm_unlink()`

PRODUCER

```
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
int main() {
    const int SIZE = 4096;
    const char *name = "OS";
    const char *message_0 = "Hello";
    const char *message_1 = "World!";

    int shm_fd;

    void *ptr;

    shm_fd = shm_open(name, O_CREAT | O_RDWR, 0666);           //shm_open usage

    ftruncate(shm_fd, SIZE);                                     //ftruncate usage

    ptr = mmap(0, SIZE, PROT_WRITE, MAP_SHARED, shm_fd, 0);    //mmap usage

    sprintf(ptr, "%s", message_0);
    ptr += strlen(message_0);
    sprintf(ptr, "%s", message_1);
    ptr += strlen(message_1);
    return 0;
}
```

CONSUMER

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <sys/shm.h>
#include <sys/stat.h>
int main() {
    const int SIZE = 4096;
    const char *name = "OS";
    int shm_fd;
    void *ptr;

    shm_fd = shm_open(name, O_RDONLY, 0666);           //shm_open usage

    ptr = mmap(0, SIZE, PROT_READ, MAP_SHARED, shm_fd, 0); //mmap usage

    printf("%s", (char *)ptr);

    shm_unlink(name);                                   //shm_unlink usage
    return 0;
}
```

MACH — MESSAGE PASSING

- Processes communicate via *messages* passed through *ports* (*buffers*).
- API includes:
 - `msg_send()`
 - `msg_receive()`
 - `msg_rpc()`
 - `port_allocate()`
 - Messages are stored FIFO but per sender.
- If ports are named, they can support n-1, 1-n, n-n communication. (...but not *broadcasts*...)


header			data		
src	dst	size	var1	var2	var3



COMMUNICATION IN CLIENT-SERVER SYSTEMS (.6)

SOCKETS

- One concept for the endpoint of a communication channel between systems is a socket.
- A socket is composed of a system's id (i.e., IP) and a port number. Ex: 98.175.148.150:80.
 - Specific ports below 1024 have semantics attached, e.g., 80 is HTTP.
- On the next slide, we'll take a look at establishing sockets and using them for communication in Java. (C is much more messy.)
 - The example will use the TCP protocol, but is another common network protocol called UDP.



USING SOCKETS IN JAVA

```
import java.net.*;
import java.io.*;
public class DateServer {
    public static void main(String[] args) {
        try {
            ServerSocket sock = new ServerSocket(6013);

            while (true) {
                Socket client = sock.accept();
                PrintWriter pout = new
                    PrintWriter(client.getOutputStream(), true);

                pout.println(new java.util.Date().toString());

                client.close();
            }
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

```
import java.net.*;
import java.io.*;
public class DateClient {
    public static void main(String[] args) {
        try {
            Socket sock = new Socket("127.0.0.1", 6013);
            InputStream in = sock.getInputStream();
            BufferedReader bin = new
                BufferedReader(new InputStreamReader(in));

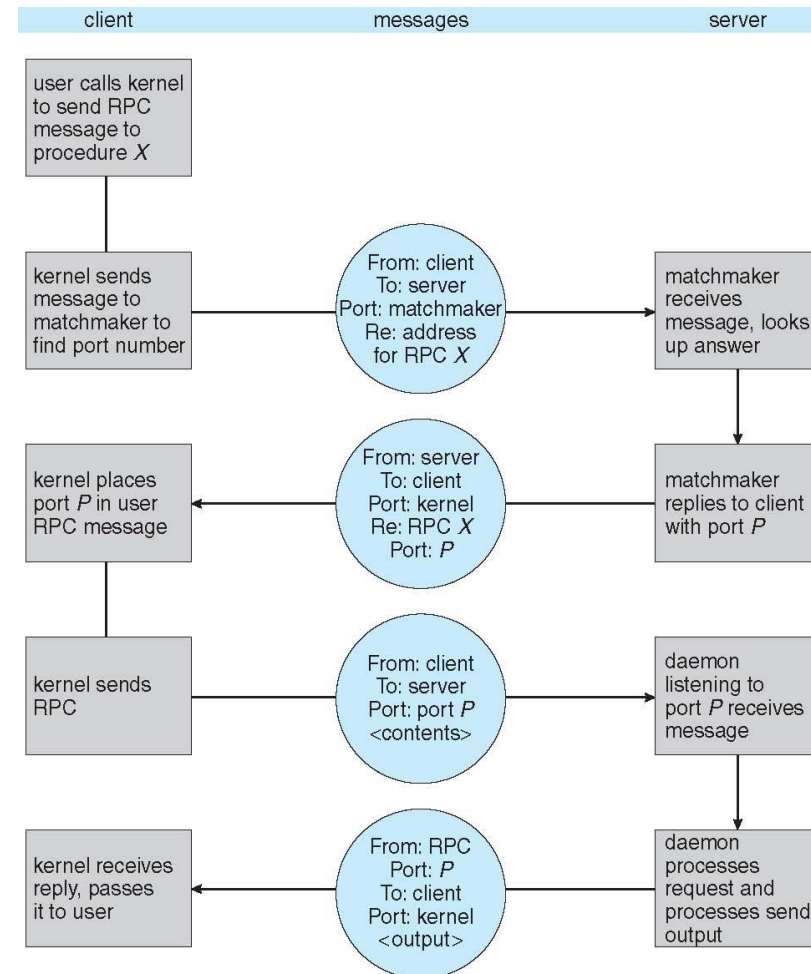
            String line;
            while ( (line = bin.readLine()) != null)
                System.out.println(line);

            sock.close();
        }
        catch (IOException ioe) {
            System.err.println(ioe);
        }
    }
}
```

“localhost”


REMOTE PROCEDURE CALLS

- The sockets example we just did is low level – we moved a stream of characters (bytes) across processes.
- A higher level mechanism is Remote Procedure Calls (RPCs).
 - Sends a well-structured packet which contains procedure identifier and input.
- Provided *API stubs* exist to *marshal* parameters into RPC format, clients act as if they are simply calling a function.
- Since calls may have some impact on the target system (i.e., they don't just return a value), we need to be sure they execute exactly once.
 - How?
- Also: RPCs are linked to specific ports so we need something like a *matchmaker* to look up the server's port.



PIPES

- A pipe is a UNIX concept for passing information between processes.
- Some considerations:
 - Bidirectional or unidirectional?
 - Half duplex or full duplex?
 - Are processes asymmetric?
 - Networked or local?
- There are two types of pipes:
 - *Ordinary*: only created in the context of a single process.
 - Have *read-end* and *write-end*. No seek ability.
 - *Named*: explicitly named resource that can be accessed by multiple processes and is persistent. Similar to a file.



PIPES

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>
#define BUFFER_SIZE 25
#define READ_END 0
#define WRITE_END 1
int main(void) {
    char write_msg[BUFFER_SIZE] = "Greetings";
    char read_msg[BUFFER_SIZE];
    int fd[2];
    pid_t pid;

    if (pipe(fd) == -1) {
        fprintf(stderr, "Pipe failed");
        return 1;
    }
    pid = fork();
    if (pid > 0) { /* parent process */
        close(fd[READ_END]);
        write(fd[WRITE_END], write_msg, strlen(write_msg) + 1);
        close(fd[WRITE_END]);
    }
    else { /* child process */
        close(fd[WRITE_END]);
        read(fd[READ_END], read_msg, BUFFER_SIZE);
        printf("read %s", read_msg);
        close(fd[READ_END]);
    }
    return 0;
}
```